

Keyush Shah

+1 (267) 278-1269 | keyush.shah12@gmail.com |



[LinkedIn](#) |



[Website](#) |



[Blog](#)

EDUCATION

University of Pennsylvania (School of Engineering and Applied Science)

Master of Science in Engineering (MSE) in Data Science; GPA 3.9/4.0

Coursework: Machine Learning, Databases, Computer Systems, Deep Learning, Generative Modeling, NLP, Machine Perception

Philadelphia, PA

Aug 2023 – May 2025

University of Delhi

Bachelor of Science in Statistics; GPA 3.87/4.0

Delhi, India

Aug 2017 – May 2020

EXPERIENCE

Senior Software Engineer, [Omnicom](#) (NYC, New York)

Aug 2025 - Present

▪ Synthetic Audiences:

1. Architected a RAG-augmented product using **Next.js + FastAPI on Vercel/Railway** enabling pharma clients to simulate patient focus groups - interviewing AI-grounded personas & **reducing 2-4 weeks of research to minutes**.
2. Developed a Go microservice fanning out **~100 concurrent LLM calls/request with bounded goroutines and context deadlines** over a JSON API decoupling serverless routes for ads/creative testing & better failure control.

▪ Brand Intelligence Compiler (Knowledge Base):

1. Engineered a Claude based multi-step agent that orchestrated **routing->retrieval->tool-use** over a pharma knowledge base enforcing a **cited/abstain contract** + a self-verifying agent that lifted extraction by **~60%**.
2. Designed an offline **eval harness** monitoring a north star GRR + diagnostic LLM judge flagging quality regressions.
3. Deployed **request-level telemetry** with async quality scoring, enabling label-free monitoring in production traffic.

▪ Led end-to-end migration of a multi-component ML Pipeline from AzureML to Databricks by re-architecting LLM Integration, data storage and maintaining orchestration in Spark Job DAGs unifying pipeline governance.

▪ Reduced CI/CD pipelines runtime by **12x** by benchmarking Docker containers & migrating to **UV for backend services**.

AI Researcher, [Penn Medicine - Computational Social Listening Lab](#) (PA, USA)

Dec 2024 – Mar 2026

IH Risk Model ((Part Time Research) [CODE LINK]

- Reduced **model tuning time by 10x** by developing an **AutoML** pipeline using sklearn for **hernia prediction (recall: 0.9)**
- Deployed GPT-4 few-shot pipeline on 10k+ redacted surgical notes, **reducing 30% noise vs. BERT** for feature labeling.
- Built a fault tolerant multi-agent orchestration system using Claude Agent SDK featuring **LLM-as-judge** observability, automated retry logic and **FastAPI + WebSocket web interface** slashing pipeline runtime from **45 to 10 min**.
- Integrated **vLLM** & **LangChain** for scalable local LLM inference deployed on A100 GPU and **modularized the backend** (OpenAI, Mistral) using **parallelized batch inference** and **memory-aware chunking** across variable length notes.

Data Science Intern (Full-time), [Universal Media](#) (PA, USA)

May 2024 – Aug 2024

- Led the development of **3+ data pipelines** using Azure Data Factory, facilitating seamless ingesting into Azure
- Drove product insights by building **Mixed Media time series models in Azure Synapse**, analyzing marketing channel impacts on media diversity metrics. Built **Power BI** dashboards to deliver actionable insights for optimizing client strategies.

Assistant Manager (Full-time), [IIFL Finance Ltd](#)

Apr 2022 – July 2023

- **Digital Adoption:** Led a product initiative to identify digitally savvy customers – engineered features & built ADF pipelines to track campaign behavior. Trained an ML model in AzureML and exposed as a REST endpoint driving disbursement by ~50%.

SELECTED PROJECTS

▪ Particle (2024):

- Built a context-aware chatbot for product and order support using FastAPI, LangChain, and Pinecone, implementing **RAG for product queries and tool-based logic with LLM fallback** for transactions.

- Engineered multi-turn memory and entity tracking, **enabling >95% accurate follow-up handling**. Automated ingestion of synthetic data, & delivered **React + Vite frontend, improving support efficiency by 40%** and user satisfaction. [\[Link\]](#)

- **Multithreaded Image Processing (2025):** Engineered a parallelized box blur algorithm in C++ using POSIX threads, achieving a 2.8x speedup (3251 ms → 1165 ms) on 4 cores by optimizing memory access, leveraging shared-memory synchronization, and **partitioning workloads across non-overlapping thread-local regions**. [\[Link\]](#)

TECHNICAL SKILLS

Programming Languages: Python, Go, C/C++, SQL, JavaScript, Typescript, R programming

ML Libraries/Frameworks: PyTorch, scikit-learn, sentence-transformers, XGBoost, HuggingFace, spaCy, OpenCV, MLflow, A/B Testing

Databases/Web Frameworks: MySQL, PostgreSQL, SSMS, MongoDB, Neo4j, React, NodeJS, Django, Flask, FastAPI

Cloud/Big Data Orchestration: AWS (S3, Glue), Azure (Data Factory, Synapse), Vercel/Railway, Kafka, Airflow, DBT, PySpark

Tools/DevOps: DataBricks, Apache Spark, Docker, Jenkins, Git, Kubernetes, pytest, CI/CD