Keyush Shah +1 (267) 278-1269| keyush.shah12@gmail.com| LinkedIn | Website| Blog

EDUCATION

University of Pennsylvania (School of Engineering and Applied Science) **Philadelphia, PA**
Master of Science in Engineering (MSE) in Data Science; ***GPA 3.9/4.0*** ***Aug 2023 – May 2025***
Coursework: Databases/Big Data Analytics, Machine Learning, NLP, AI, Computer Vision, Deep Learning, Statistical Modeling

Kirori Mal College (KMC – University of Delhi) **Delhi, India**
Bachelor of Science (BS) in Statistics; **GPA: 8.11/10** **July 2017 – Aug 2020**

EXPERIENCE

ML Researcher, ***Computational Social Listening Lab (UPenn)*** ***May 2024 – Present***

Project 1: (**Misinformation**)

- Worked on identifying misinformation on social media and whether it relates to health outcomes segmented by race.
- Extracted linguistic features from posts using DLATK and applied LDA for topic modeling and performed correlation analysis.
- Built NLP pipelines to detect health-related misinformation from ~20K social media posts using RoBERTa and LDA topic modeling; improved **classification precision by 20%** through alignment-based entailment.
- Unified data from various online survey platforms (**10K+ responses**) in a secured server via MySQL and performed feature engineering in Pandas to prepare data for further downstream tasks.

Project 2: (**IH Risk Model**)

- Developed a pipeline to predict Incisional Hernia (IH) risk from unstructured operative notes and structured EHR data in a surgical cohort of 10k+ patients.
- Engineered a few-shot GPT-based extraction pipeline to label operative features (e.g., incision, ostomy) from redacted notes **and reduced noisy extractions by 30% vs. BERT embeddings.**

Data Science Intern (Full-time)**,** ***Universal Media (PA, USA)*** ***May 2024 – Aug 2024***

- Architected Azure SQL Database solutions, encompassing DDL scripts to enhance data management and reporting solutions.
- Led the development of **3+ data pipelines** using Azure Data Factory (ADF), facilitating the seamless ingestion and transformation of diverse data sources into the Azure environment.
- Developed python scripts for data transformation, stored them in Blob storage and executed them via batch activity in ADF.
- Drove product marketing insights by building **Mixed Media models (MMM) in Azure Synapse**, analyzing marketing channel impacts on media diversity metrics. Built **Power BI** dashboards to deliver actionable insights for optimizing client strategies.
- Authored **5 stored procedures** in SQL, automating repetitive tasks and improving query performance by over 30%.

Assistant Manager (Full-time), ***IIFL Finance Ltd*** ***Apr 2022 – July 2023***

- Analyzed ETL process failures and created **10+** paginated reports in SSIS to help the management track 1000+ branches.
- Optimized & migrated complex SQL queries from an obsolete database server that improved the **reporting services by ~40%.**
- **Digital Adoption:** Led a product-focused initiative to identify digitally savvy customers by engineering features and building ADF pipelines to track campaign behavior. Trained and deployed a Random Forest model (**with a 90% accuracy**) in Azure ML Studio; **exposed it as a REST endpoint consumed by marketing campaigns,** driving digital **disbursal adoption by 50%.**

SELECTED PROJECTS

- **Azure ETL (2025):** Built a scalable ETL pipeline in Azure Data Factory that was capable of ingesting and transforming 1M+ rows daily. Ingested raw data into Data Lake Storage Gen2, achieved **~50% reduction in query latency** via Azure Databricks optimizations, and analyzed it in Synapse Analytics - delivering a seamless end-to-end solution. [Link]
- **FitBit(2024):** Engineered a Django health chatbot leveraging PostgreSQL for robust patient data management, featuring an LLM-agnostic architecture with seamless model switching via **Langchain that reduced overhead by 40%**. Optimized memory usage for handling long conversations and implemented entity extraction to dynamically enhance medical context. [Link]
- **Diffusion Transformer (2024):** Implemented PatchVAE with convolutional encoders and patch-based decoding for fine-grained feature extraction. Trained a Diffusion Transformer to sample from the latent space of PatchVAE, achieving a **30% reduction** in FID score and **2x greater feature diversity** compared to VAE-generated samples. [Link]
- **Statistical Segmentation (2024):** Modeled customer purchase behavior using Pareto II and Weibull distributions; optimized parameters using Excel Solver and back tested on historical unit sales, improving segment alignment with **actual revenue by 21%** and informing future campaign targeting. (Solver, Marketing). [Link]

TECHNICAL SKILLS

Programming Languages: Python, SQL, C/C++, R programming, JavaScript
Databases/Frameworks: MySQL, PostgreSQL, SSMS, MongoDB, Neo4j, PyTorch, React, NodeJS, A/B Testing, PySpark, Django
Cloud/Big Data Orchestration: AWS (S3, Glue), Azure (Data Factory, Synapse, DevOps), Docker, DataBricks, Kafka, Airflow, DBT